

Introduction to Cypher

Table of Contents

About this module	1
What is Cypher?	1
Nodes	2
Labels	2
Examining the schema	3
Using MATCH to retrieve nodes	3
Exercise 1: Retrieving nodes.	5
Properties	5
Examining property keys	6
Retrieving nodes filtered by a property value	7
Returning property values	8
Exercise 2: Filtering queries using property values	9
Relationships	9
ASCII art	10
Querying using relationships	10
Examining relationships	10
Using a relationship in a query	11
Using anonymous nodes in a query	13
Using an anonymous relationship for a query	14
Retrieving the relationship types	15
Retrieving properties for relationships	16
Using patterns for queries	17
Cypher style recommendations	20
Exercise 3: Filtering queries using relationships	21
Check your understanding	22
Question 1	22
Question 2	22
Question 3	22
Summary	23

About this module

Cypher is the query language you use to retrieve data from the Neo4j Database, as well as create and update the data. In this module, **you will learn** how to write Cypher **code** to retrieve data as nodes in a graph, as well as property values for nodes retrieved. **You will learn** how to filter queries based upon the **Labels** and property values for nodes. **You will also learn** how to specify queries where the relationships between nodes are used to retrieve data.

At the end of this module, you be able to write Cypher **code** to:

- Retrieve nodes from the graph.
- Filter nodes retrieved using **property values of nodes.**
- Retrieve property values from nodes in the graph.
- Filter nodes retrieved using relationships.

Throughout this training, you should refer to:

- [Neo4j Developer Manual](#)
- [Cypher Reference card](#)

What is Cypher?

Cypher is a declarative query language that allows for expressive and efficient querying and updating of graph data. Cypher is a relatively simple and very powerful language. Complex database queries can easily be expressed through Cypher, allowing you to focus on your domain instead of getting lost in database access.

Cypher is designed to be a human-friendly query language, suitable for both developers and **any** other professionals. **It's** guiding goal is to make the simple things easy, and the complex things possible. Optimized for being read by humans, Cypher's construct uses English prose and iconography (called ASCII Art) to make queries more self-explanatory.

Being a declarative language, Cypher focuses on the clarity of expressing **what** to retrieve from a graph, not on **how** to retrieve it. This is in contrast to imperative, programmatic APIs for database access. This approach makes query optimization an implementation detail instead of a burden on the developer, removing the requirement to update all traversals just because the physical database structure has changed.

Cypher is inspired by a number of different approaches and builds upon established practices for expressive querying. Many of the Cypher keywords like **V** and **ORDER BY** are inspired by SQL. The pattern matching functionality of Cypher borrows **expression approaches** from SPARQL. And some of the collection semantics have been borrowed from languages such as Haskell and Python.

The Cypher language has been made available to anyone to implement and use **openCypher** (openCypher.org), allowing any database vendor, researcher or other interested party to reap the benefits of our years of effort and experience in developing a first class graph query language.

Nodes

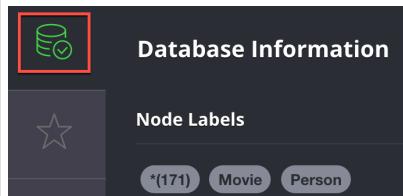


Cypher uses a pair of parentheses like `()`, `(n)` to represent a node. Recall that a node typically represents an entity in your domain. An anonymous node, `()`, represents one or more nodes during a query processing where there are no restrictions of the type of node or the properties of the node. When you specify `(n)` for a node, you are telling the query processor that for this query, use the variable *n* to represent nodes that will be processed later in the query for further query processing or for returning values from the query.

Labels

Nodes in a graph are typically labeled. Labels are used to group nodes and filter queries against the graph. That is, labels can be used to optimize queries. In the *Movie* database you will be working with, the nodes in this graph are labeled *Movie* or *Person* to represent two types of nodes.

For example, you can see the labels in the database by simply clicking the Database icon in Neo4j Browser:



You can filter the types of nodes that you are querying, by specifying a **label** for a node. A node can have zero or more labels. Labels are useful if you want to filter the types of nodes you are retrieving from the graph.

Here is the simplified syntax for specifying a node:

```
( ) |  
('variable') |  
(:'Label') |  
('variable':'Label') |  
(:'Label1':'Label2') |  
('variable':'Label1':'Label2')
```

Notice that a node must have the parentheses. The labels and the variable for a node are optional.

Here are examples of specifying nodes in Cypher:

```
( ) // anonymous node that will not be referenced later in the query  
processing  
(p) // variable p holds a reference to a node that will be used later  
(:Person) // anonymous node of type Person  
(p:Person) // p holds a reference to a node of type Person  
(p:Actor:Director) // p holds a reference to a node of both types Actor and Director
```

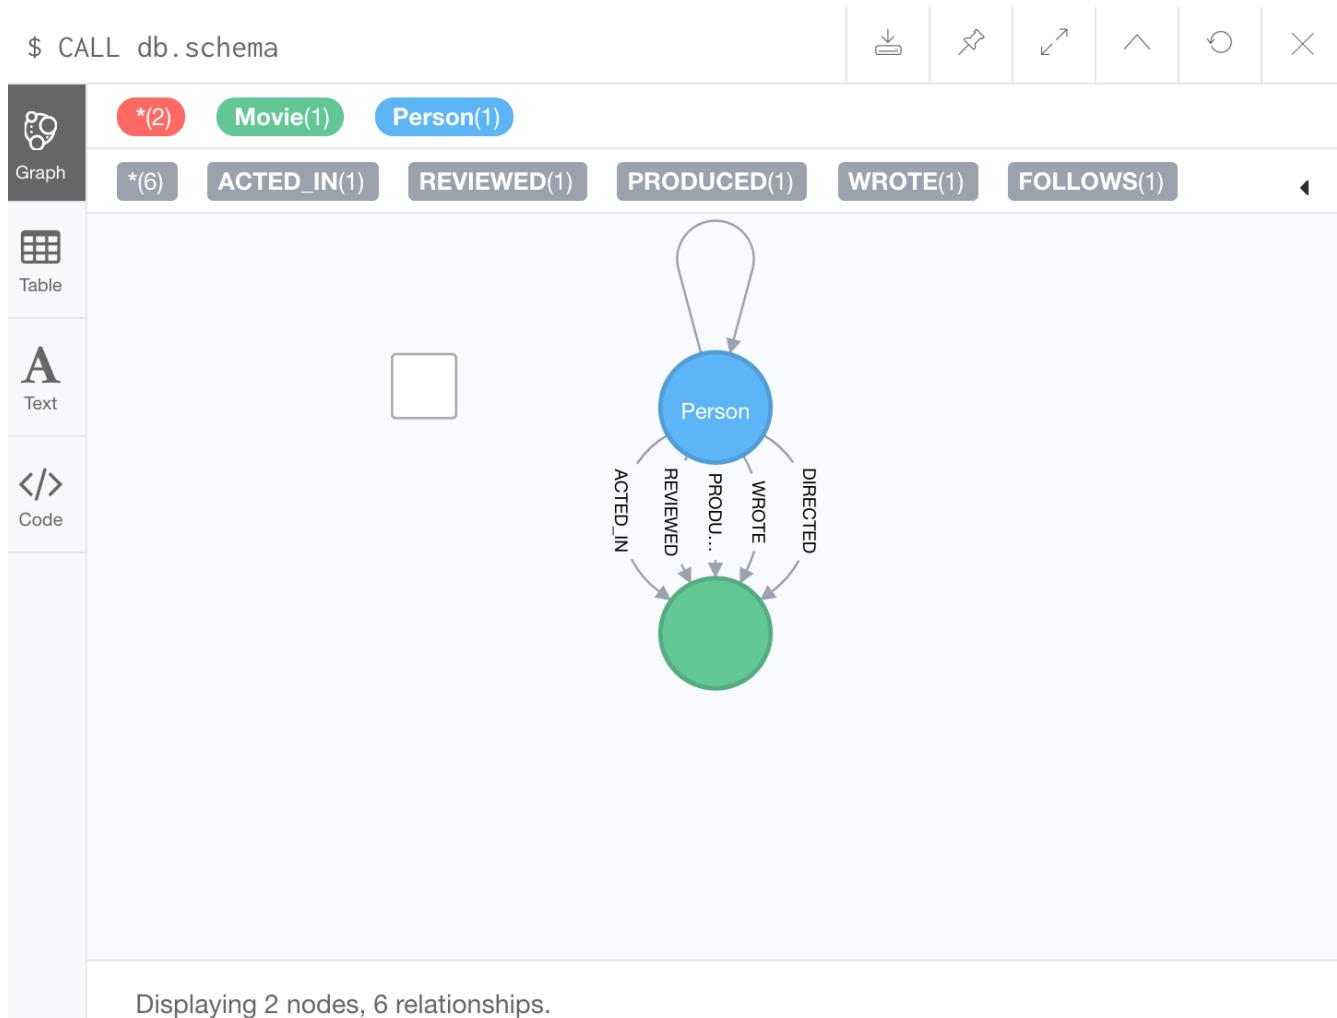
A node can have multiple labels. For example a node can be created with a label of *Person* and that same node can be modified to also have the label of *Actor* and/or *Director*. In Cypher you can place a comment (starts with `//`) anywhere in your Cypher code to specify that the rest of the line is

interpreted as a comment.

Examining the schema

When you are first learning about the data (nodes, labels, etc.) in a graph, it is helpful to examine the schema of the graph. You do so by executing `CALL db.schema`, which calls the Neo4j ~~HTTP~~ method that returns information about the nodes, labels, and relationships in the graph.

For example, when we run this ~~method~~ in our training environment, we see the following in the result pane. Here we see that the graph has 2 labels defined for nodes, *Person* and *Movie*. Each type of nodes is displayed in a different color. The relationships between nodes are also displayed, which you will learn about later in this module.



Using MATCH to retrieve nodes

The most widely-used Cypher clause is `MATCH`. The `MATCH` clause performs a query against the graph based upon the criteria for the query. During the query processing, the graph engine traverses the graph to find all nodes that match the quer ~~teria~~ teria. As part of query, you can return nodes or data from the nodes using the `RETURN` clause. The `RETURN` clause must be the last clause of a query to the graph. Later in this training, you will learn how to use `MATCH` to select nodes and data for updating the graph. First, you will learn how to simply return nodes.

Here is the simplified syntax for a query:

```
MATCH ('variable') RETURN 'variable' |  
MATCH ('variable':'Label') RETURN 'variable'
```

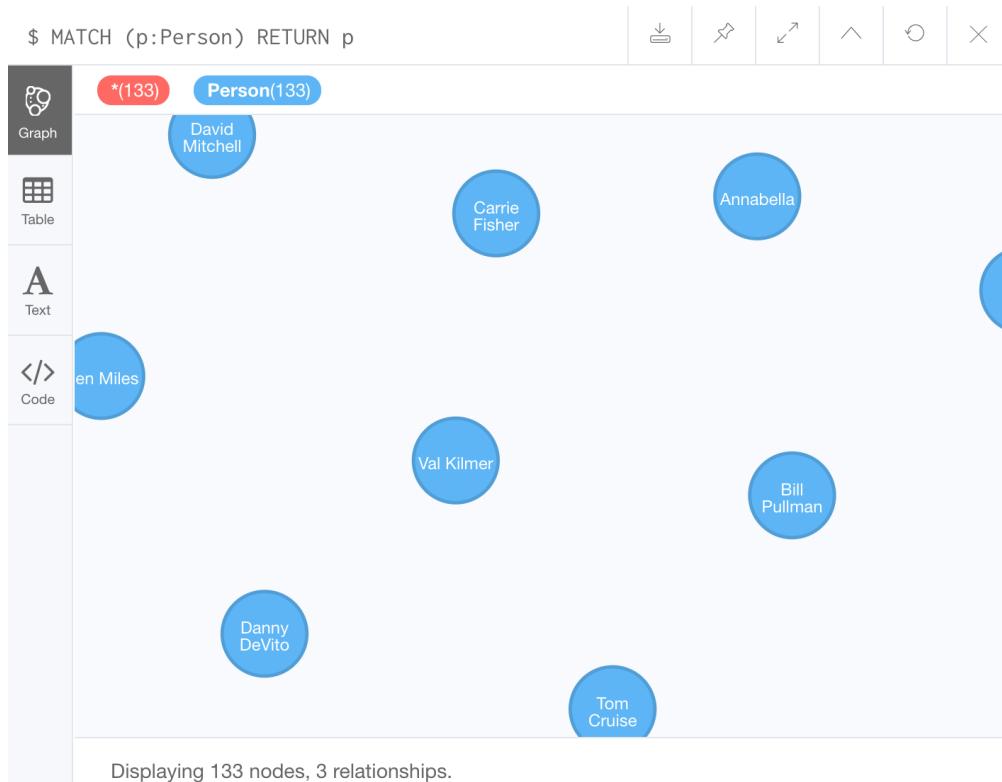
The **MATCH** clause contains the query and has a **RETURN** clause that [] ed to return values retrieved. Notice that the Cypher keywords **MATCH** and **RETURN** are upper-case. This coding convention is a [] practice and will be followed through this training. This **MATCH** clause returns all nodes a [] their [] relationships in the graph, where the optional *Label* is used to filter the query. The *variable* must be specified here, otherwise the query will have nothing to return.

Here are example queries to the *Movie* database:

```
MATCH (n) RETURN n      // returns graph of all nodes  
MATCH (p:Person) RETURN p // returns graph of all Person nodes
```

When we execute the Cypher code, **MATCH (p:Person) RETURN p**, the graph engine returns all nodes with the label *Person*. The default view of the returned nodes are the nodes that were referenced by the variable *p*.

The result returned is:



We can also view the nodes as table data where the nodes and their associated property values are shown in a JSON-style format:

```
$ MATCH (p:Person) RETURN p
```



The screenshot shows the Neo4j Browser interface with the 'Table' tab selected. On the left, there's a sidebar with icons for Graph, Table (selected), Text, and Code. The main area displays three nodes, each represented by a gray box containing JSON-like data:

- Node 1: { "name": "Keanu Reeves", "born": 1964 }
- Node 2: { "name": "Carrie-Anne Moss", "born": 1967 }
- Node 3: { "name": "Laurence Fishburne", "born": 1961 }

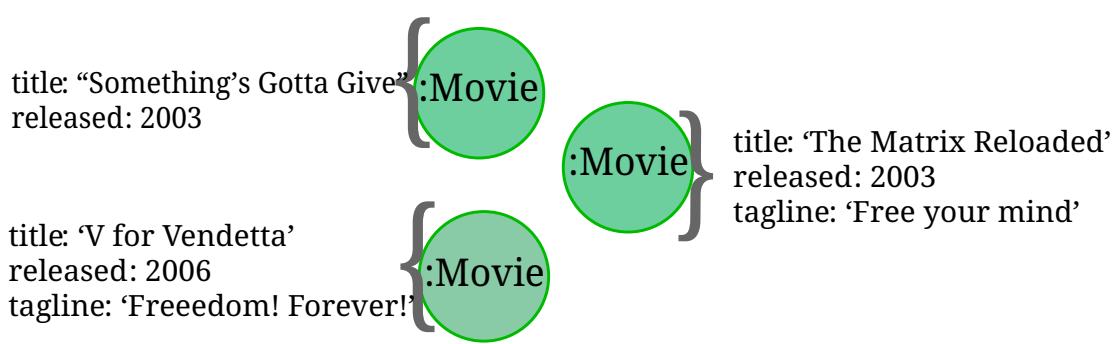
Exercise 1: Retrieving nodes

In the query edit pane of Neo4j Browser, execute the browser command: `:play http://guides.neo4j.com/intro-neo4j-exercises` and follow the instructions for Exercise 1.

Properties

In Neo4j, a node (and a relationship, which you will learn about later) can have properties that are used to further define a node. A property is identified by its property key. Recall that nodes are used to represent the entities of your business model. A property is defined for a node and not for a type of node. All nodes of the same type need not have the same properties.

For example, in the *Movie* graph, all *Movie* nodes have both *title* and *released* properties. However, it is not a requirement that every *Movie* node has a property, *tagline*.



Properties can be used to filter queries so that a subset of the graph is retrieved. In addition, with the `RETURN` clause, you can return `v`'s from the retrieved nodes, rather than the nodes.

Examining property keys

As you prepare to create Cypher queries that use property values to filter a query, you can view the values for property keys of a graph by simply clicking the Database icon in the Neo4j Browser. Alternatively, you can execute `CALL db.propertyKeys`, which calls the ~~Neo4j library~~ method that returns the property keys for the graph.

Here is what you will see in the result pane when you call the method to return the property keys in the *Movie* graph. This result stream contains all property keys in the graph. It does not display which nodes utilize these property keys.

\$ CALL db.propertyKeys	
Table Text Code	<pre>propertyKey "title" "released" >tagline" "name" "born" "roles" "summary" "rating" "id" "share_link" "favorite_count" "display_name"</pre>
	Started streaming 12 records in less than 1 ms and completed in less than 1 ms.

Retrieving nodes filtered by a property value

You have learned previously that you can filter node retrieval by specifying a label. Another way you can filter a retrieval is to specify a value for a property. Any node that matches the value will be retrieved.

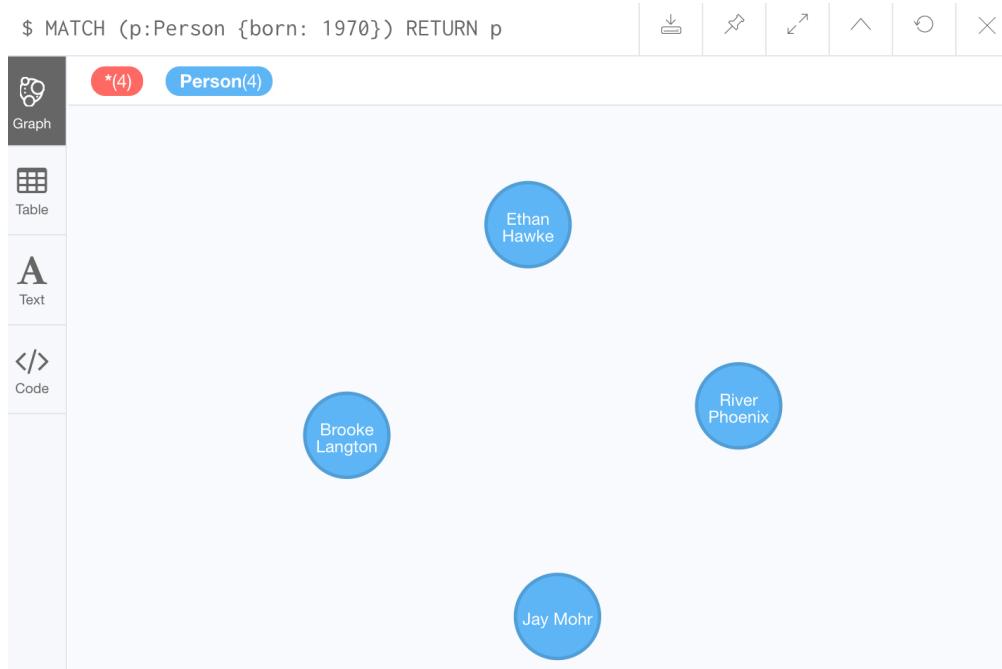
Here is the simplified syntax for a query where we specify one or more values for properties that will be used to filter the query results and return a subset of the graph:

```
MATCH ('variable' {'propertyKey': 'PropertyValue'}) RETURN 'variable' |  
MATCH ('variable':'Label' {'propertyKey': 'PropertyValue'}) RETURN 'vari |  
MATCH ('variable' {'propertyKey1': 'PropertyValue1', 'propertyKey2':  
'PropertyValue2'}) RETURN 'variable' |  
MATCH ('variable':'Label' {'propertyKey': 'PropertyValue' , 'propertyKey2':  
'PropertyValue2'}) RETURN 'variable' |
```

Here is an example where we filter the query results using a property value. We only retrieve Person nodes that have a *born* property value of 1970.

```
MATCH (p:Person {born: 1970}) RETURN p
```

The result returned is:



Here is an example where we specify two property values for the query.

```
MATCH (m:Movie {released: 2003, tagline: "Free your mind"}) RETURN m
```

Here is the result returned:

```
$ MATCH (m:Movie {released: 2003, tagline: "Free your mind"}) RETURN m
```



m

```
{  
    "title": "The Matrix Reloaded",  
    "tagline": "Free your mind",  
    "released": 2003  
}
```

Started streaming 1 records after 1 ms and completed after 2 ms.

As it turns out, there is only one movie with the *tagline*, 'Free your mind' in the *Movie* database, but if we had specified a different year, the query would not have returned a value because when you specify properties, both properties must match.

Returning property values

Thus far, you have seen how to retrieve nodes and return nodes (entire graph or a subset of the graph). You can use the **RETURN** clause to return property values of nodes retrieved.

Here is the simplified syntax for returning property values, rather than nodes.

```
MATCH ('variable' {'propertyKeySpec1'}) RETURN 'variable'.propertyKey2' |  
MATCH ('variable'[:'Label'] {'propertyKeySpec1'}) RETURN 'variable'.propertyKey2' |  
MATCH ('variable' {'propertyKeySpec11'}) RETURN 'variable'.propertyKey2',  
'variable'.propertyKey3' |  
MATCH ('variable'[:'Label'] {'propertyKeySpec1'}) RETURN 'variable'.propertyKey2',  
'variable'.propertyKey3'
```

In this example, we use the *born* property to refine the query, but rather than returning the nodes, we return the *name* and *born* values for every node that satisfies the query.

```
MATCH (p:Person {born: 1965}) RETURN p.name, p.born
```

The result returned is:

```
$ MATCH (p:Person {born: 1965}) RETURN p.name, p.born
```



Table

A
Text

</>
Code

p.name	p.born
"Lana Wachowski"	1965
"Tom Tykwer"	1965
"John C. Reilly"	1965

Started streaming 3 records in less than 1 ms and completed after 1 ms.

If you want to customize the headings for a table containing property values, you can specify aliases for property values.

```
MATCH ('variable':'Label' {'propertyKey1': 'PropertyValue1'}) RETURN  
'variable'.propertyKey2 AS 'alias2'
```

Note: If you want the heading to contain a space between strings, you must specify the alias with the back tick `` character, rather than a single or double quote character.

Here we specify aliases for the returned property values:

```
MATCH (p:Person {born: 1965}) RETURN p.name AS Name, p.born AS 'Birth Year'
```

The result returned is:

```
$ MATCH (p:Person {born: 1965}) RETURN p.name AS Name, p.born AS BirthYear
```



Table

A
Text

</>
Code

Name	BirthYear
"Lana Wachowski"	1965
"Tom Tykwer"	1965
"John C. Reilly"	1965

Exercise 2: Filtering queries using property values

In the query edit pane of Neo4j Browser, execute the browser command: :play <http://guides.neo4j.com/intro-neo4j-exercises> and follow the instructions for Exercise 2.

Relationships

The most important concept in a Neo4j graph is relationships. Relationships are what make Neo4j

graphs such a powerful tool for connecting complex and deep data. A relationship is a **directed** connection between two nodes that has a **relationship type** (+). In addition, a relationship can have properties, just like nodes. In a graph where you want to retrieve nodes, you can use relationships between nodes to filter a query.

ASCII art

Thus far, you have learned how to specify a node in a **MATCH** clause. You can specify nodes and their relationships to traverse the graph and quickly find the data of interest.

Here is how Cypher uses ASCII art to specify path used for a query:

```
()          // a node
()---() // 2 nodes have some type [ ] relationship
()-->() // the first node has a relationship to the second node
()<--() // the second node has a relationship to the first node
```

Querying using relationships

In your **MATCH** clause, you specify how you want a relationship to be used to perform the query. The relationship can be specified with or without direction.

Here is the simplified syntax for retrieving a set of nodes that satisfy one or more direct [] relationships:

```
MATCH ('node1-spec')-[REL_TYPE]->('node2-spec') RETURN 'node1-var', 'node2-var' |
MATCH ('node1-spec')-[REL_TYPEA | REL_TYPEB]->('node2-spec') RETURN 'node1-var',
'node2-var' |
```

where:

node1-spec is a specification of a node where you may include node labels and property values for filtering and uses the variable, *node1-var* during the query.

:REL_TYPE is the type (~~label~~) for the relationship. For this syntax the relationship is from *node1* to *node2*.

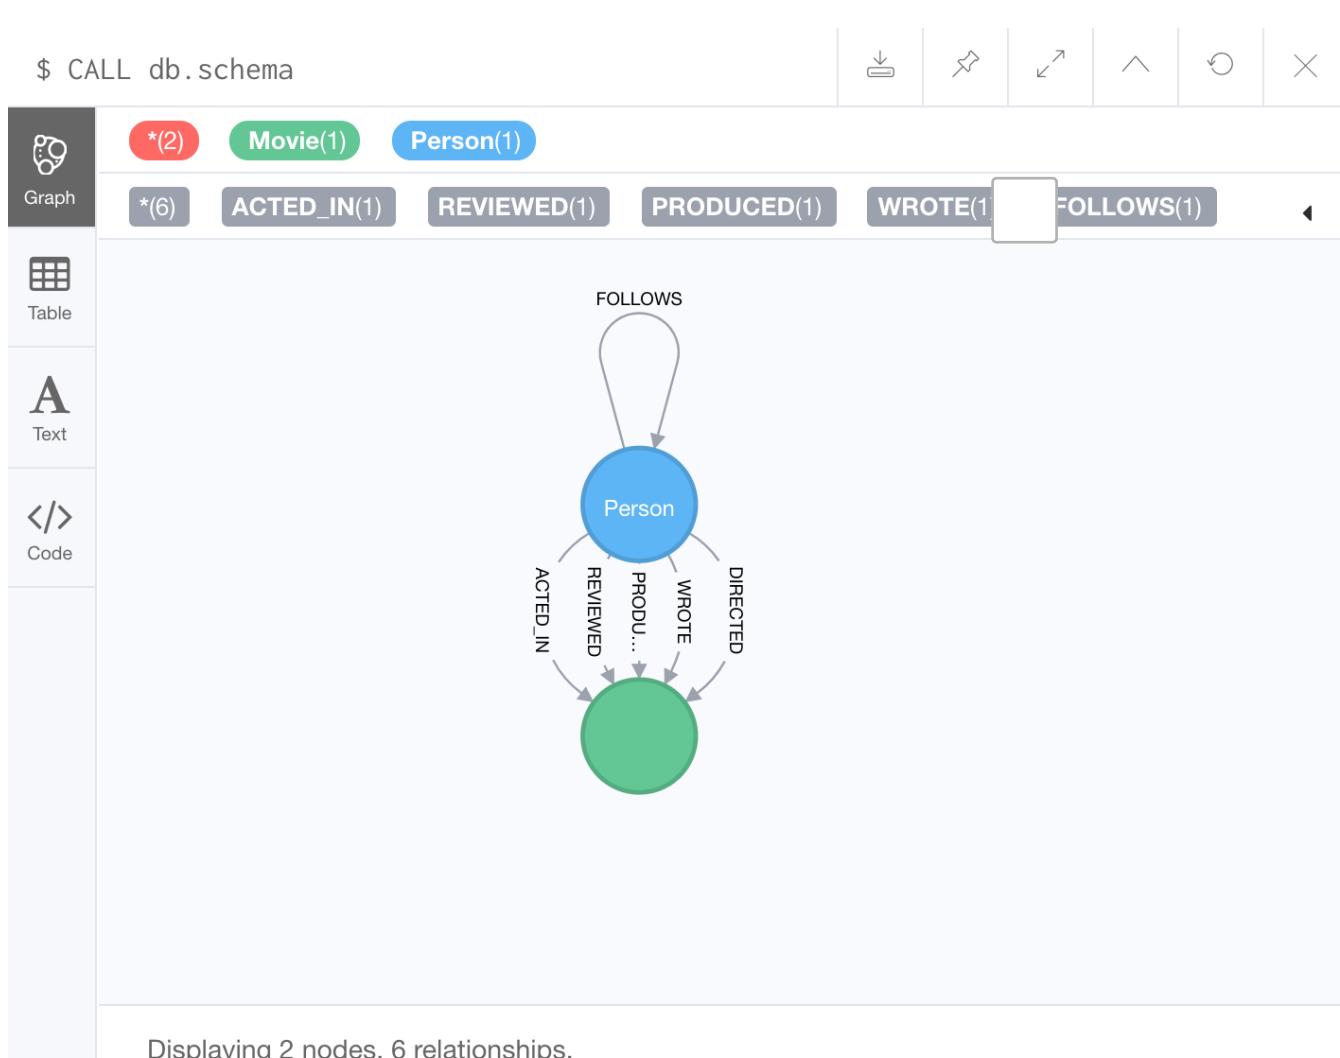
:REL_TYPEA, *:REL_TYPEB* are the relationships from *node1* to *node2*. The nodes are returned if at least one of the relationships exists.

node2-spec is a specification of a node where you may include node labels and property values for filtering and uses the variable, *node2-var* during the query.

Examining relationships

You can run **CALL db.schema** to view the relationships in the graph. In the *Movie* graph, we see these relationships between the nodes:

Here we see that this graph has a total of 6 relationships between the nodes. Some *Person* nodes are connected to other *Person* nodes using the *FOLLOW*s relationship type. All of the other relationships in this graph are from *Person* nodes to *Movie* nodes.



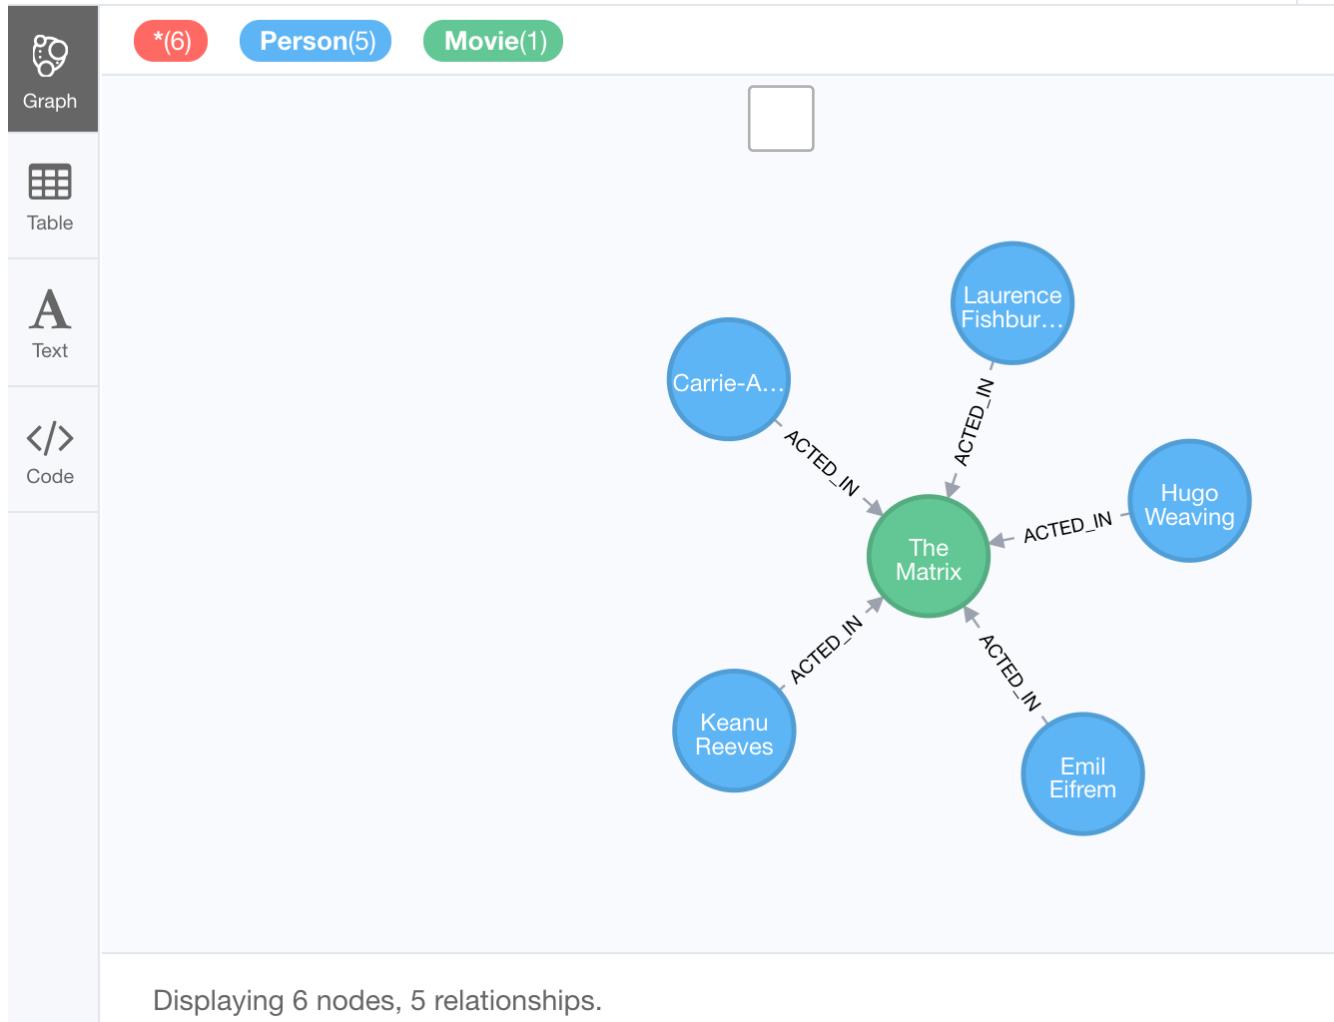
Using a relationship in a query

Here is an example where we retrieve the *Person* nodes that have the *ACTED_IN* relationship to the *Movie*, 'The Matrix'. In other words, show me the actors that acted in 'The Matrix'.

```
MATCH (p:Person)-[:ACTED_IN]->(m:Movie {title: 'The Matrix'}) RETURN p, m
```

The result returned is:

```
$ MATCH (p:Person)-[:ACTED_IN]->(m:Movie {title:'The Matrix'}) RETURN p,m
```



For this query, we are using the variable *p* to represent the *Person* nodes during the query and the variable *m* to represent the *Movie* node retrieved. We return a graph with the *Person* nodes, the *Movie* node and their ~~ACTED_IN~~ relationships.

Important: You specify node labels whenever possible in your queries as it optimizes the retrieval in the graph engine. That is, you should **not** specify this same query as ~~MATCH (p)-[:ACTED_IN]->(m:Movie {title:'The Matrix'}) RETURN p,m~~.

Here is another example where we want to know the movies that *Tom Hanks* acted in and directed:

```
MATCH (p:Person {name: 'Tom Hanks'})[:ACTED_IN | :DIRECTED]->(m:Movie)  
RETURN p.name, m.title
```

The result returned is:

```
$ MATCH (p:Person {name: 'Tom Hanks'})-[:ACTED_IN | :DIRECTED]->(m:Movie) RETURN p.name, m.title
```

p.name	m.title
"Tom Hanks"	"Apollo 13"
"Tom Hanks"	"Cast Away"
"Tom Hanks"	"The Polar Express"
"Tom Hanks"	"A League of Their Own"
"Tom Hanks"	"Charlie Wilson's War"
"Tom Hanks"	"Cloud Atlas"
"Tom Hanks"	"The Da Vinci Code"
"Tom Hanks"	"The Green Mile"
"Tom Hanks"	"You've Got Mail"
"Tom Hanks"	"That Thing You Do"
"Tom Hanks"	"That Thing You Do"
"Tom Hanks"	"Joe Versus the Volcano"
"Tom Hanks"	"Sleepless in Seattle"

Started streaming 13 records after 1 ms and completed after 1 ms.

Notice that there are multiple rows returned for the movie, *That Thing You Do*. This is because *tom Hanks* acted in and directed that movie.

Using anonymous nodes in a query

Suppose you wanted to retrieve the actors that acted in 'The Matrix', but you do not need any information returned about the *Movie* node. You need not specify a variable for a node in a query if that node is not returned or used for later processing in the query. You can simply use the anonymous node in the query as follows:

```
MATCH (p:Person)-[:ACTED_IN]->(:Movie {title: 'The Matrix'}) RETURN p.name
```

The result returned is:

```
$ MATCH (p:Person)-[:ACTED_IN]->(:Movie {title:'The Matrix'}) RETURN p.name
```



p.name

No node variable specified here

"Emil Eifrem"

"Hugo Weaving"

"Laurence Fishburne"

"Carrie-Anne Moss"

"Keanu Reeves"

Started streaming 5 records after 1 ms and completed after 2 ms.

Note: A best practice is to place named nodes (those with variables) before anonymous nodes in a **MATCH** clause.

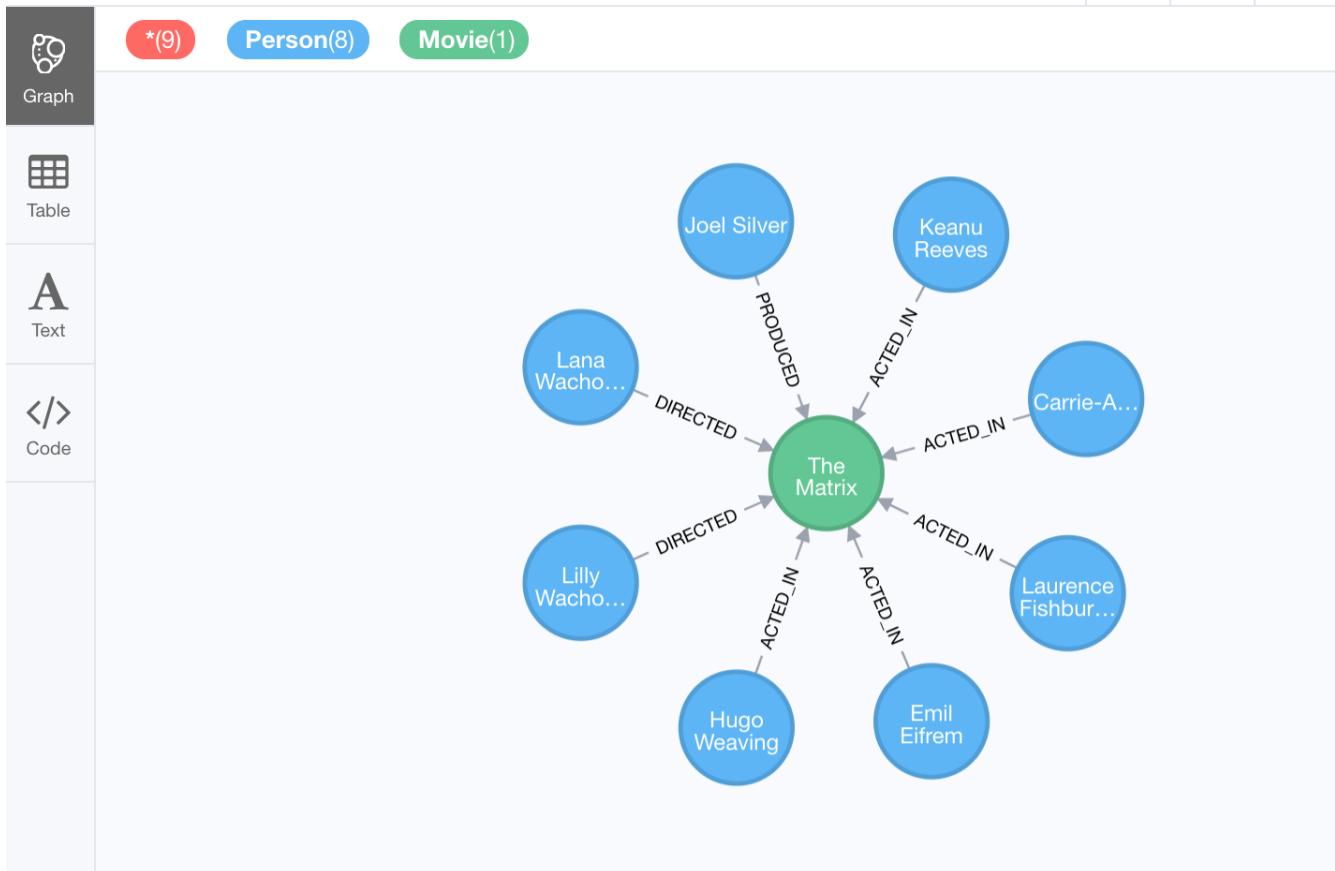
Using an anonymous relationship for a query

Suppose you want to find all people who are in any way connected to the movie, 'The Matrix'. You can specify an empty relationship type in the query so that all relationships are traversed and the appropriate results are returned. In this example, we want to retrieve all *Person* nodes that have any type of connection to the *Movie* node, with the *title* 'The Matrix'. This query returns more nodes with the relationships types, *DIRECTED*, *ACTED_IN*, and *PRODUCED*.

```
MATCH (p:Person)-->(:Movie {title: 'The Matrix'}) RETURN p, m
```

The result returned is:

```
$ MATCH (p:Person)-->(m:Movie {title: 'The Matrix'}) RETURN p, m
```



Note: You can also specify an anonymous relationship with the empty bracket For example:

MATCH (p:Person)-[]->(m:Movie {title: 'The Matrix'}) RETURN p, m

- MATCH (p:Person)-[]-(m:Movie {title: 'The Matrix'}) RETURN p, m
- MATCH (m:Movie) -[]-(p:Person {name: 'Keanu Reeves'}) RETURN p, m

In this training, we will use and to represent anonymous relationships as it is a Cypher best practice.

Retrieving the relationship types

There is a built-in function, `type()` that returns the relationship type of a relationship. Here is an example where we use the `rel` variable to hold the relationships retrieved. We then use this variable to return the relationship types.

```
MATCH (p:Person)-[rel]->(:Movie {title:'The Matrix'}) RETURN p.name, type(rel)
```

The result returned is:

```
$ MATCH (p:Person)-[rel]->(:Movie {title:'The Matrix'}) RETURN p.name, type(rel)
```

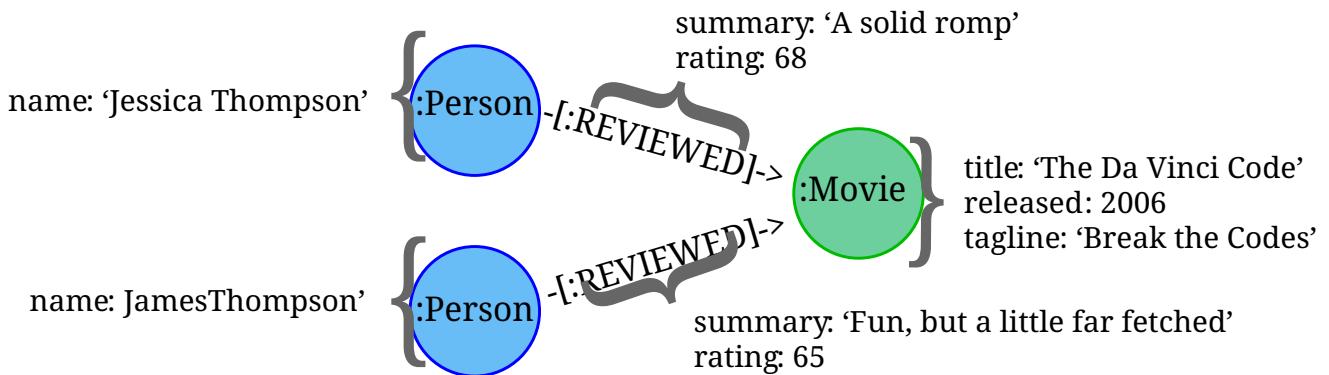
	p.name	type(rel)
Table	"Emil Eifrem"	"ACTED_IN"
A Text	"Joel Silver"	"PRODUCED"
	"Lana Wachowski"	"DIRECTED"
	"Lilly Wachowski"	"DIRECTED"
	"Hugo Weaving"	"ACTED_IN"
	"Laurence Fishburne"	"ACTED_IN"
	"Carrie-Anne Moss"	"ACTED_IN"
	"Keanu Reeves"	"ACTED_IN"

Started streaming 8 records after 1 ms and completed after 2 ms.

Retrieving properties for relationships

Recall that a node can have a set of properties, each identified by its property key. Relationships can also have properties. This enables your graph model to provide more data about the relationships between the nodes.

Here is an example from the *Movie* graph. The movie, 'The Da Vinci Code' has two people that reviewed it, Jessica Thompson and James Thompson. Each of these *Person* nodes has the *REVIEWED* relationship to the *Movie* node for 'The Da Vinci Code'. Each relationship has properties that further describe the relationship using the *summary* and *rating* properties.



Just as you can specify property values for filtering nodes for a query, you can specify property values for a relationship. This query returns the name of the person who gave the movie a rating

of 65.

```
MATCH (p:Person)-[:REVIEWED {rating: 65}]->(:Movie {title: 'The Da Vinci Code'})  
RETURN p.name
```

The result returned is:

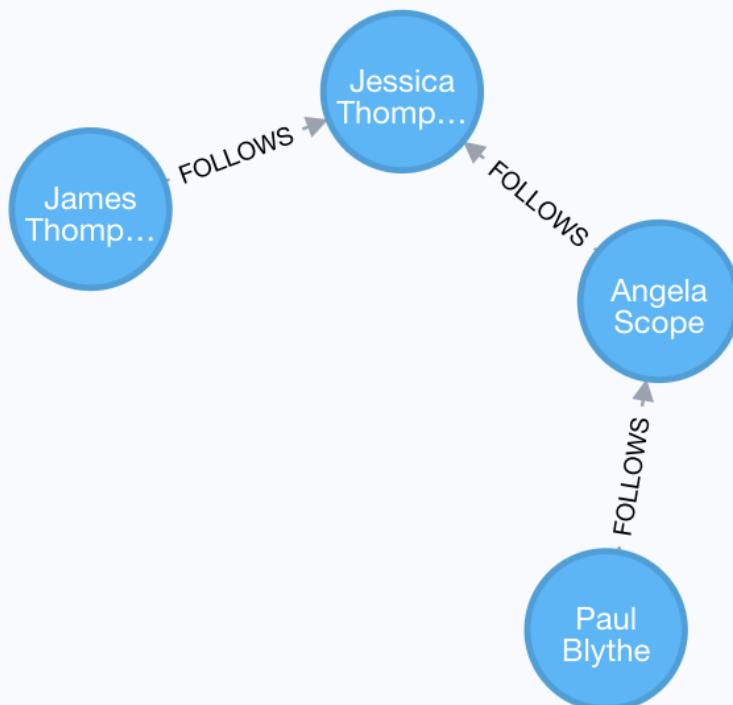
```
$ MATCH (p:Person)-[:REVIEWED {rating: 65}]->(:Movie {title: 'The Da Vinci Code'}) RETURN p.name
```

Table	p.name
A Text	"James Thompson"

Using patterns for queries

Thus far, you have learned how to specify nodes, properties, and relationships in your Cypher queries. Since relationships are directional, it is important to understand how patterns are used in graph traversal during query execution. How a graph is traversed for a query depends on what directions are defined for relationships and how the pattern is specified in the `MATCH` clause.

Here is an example of where the `FOLLOW`s relationship is used in the *Movie* graph. Notice that this relationship is directional.



We can perform a query that returns all *Person* nodes who follow 'Angela Scope':

```
MATCH (p:Person)-[:FOLLOWS]->(:Person {name:'Angela Scope'}) RETURN p
```

The result returned is:



If we reverse the direction in the pattern, the query returns different results:

```
MATCH (p:Person)<-[:FOLLOWS]-(:Person {name:'Angela Scope'}) RETURN p
```

The result returned is:

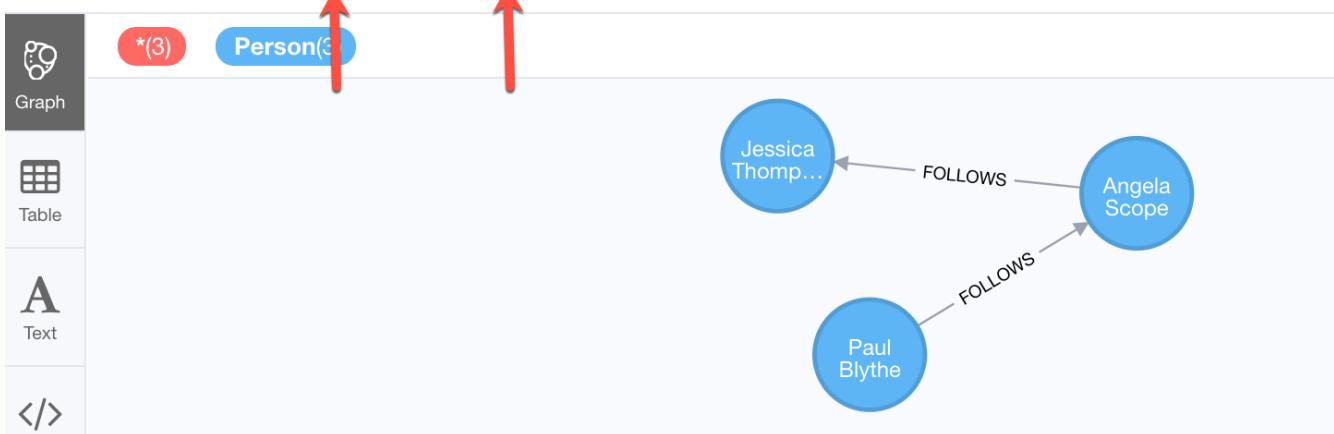


We can also find out what *Person* nodes are connected by the *FOLLOWED* relationship in either direction by removing the directional arrow from the pattern.

```
MATCH (p1:Person)-[:FOLLOWS]-(p2:Person {name:'Angela Scope'}) RETURN p1, p2
```

The result returned is:

```
$ MATCH (p1:Person)-[:FOLLOWS]-(p2:Person {name:'Angela Scope'}) RETURN p1, p2
```

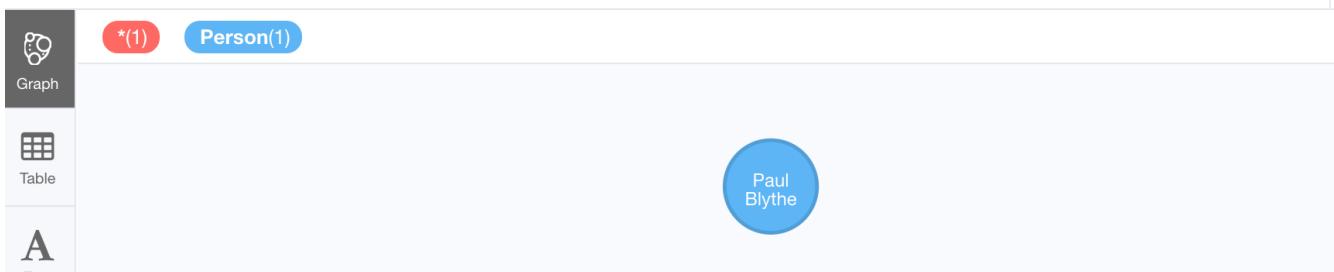


Since we have a graph, we can traverse through nodes to obtain relationships further into the traversal. For example, we can write a Cypher query to return all followers of the followers of Jessica Thompson.

```
MATCH (p:Person)-[:FOLLOWS]->(:Person)-[:FOLLOWS]->(:Person {name:'Jessica Thompson'}) RETURN p
```

The result returned is:

```
$ MATCH (p:Person)-[:FOLLOWS]->(:Person)-[:FOLLOWS]->(:Person {name:'Jessica Thompson'}) RETURN p
```



When querying the relationships in a graph, you can take advantage of the direction of the relationship to traverse the graph. For example, suppose we wanted to get a result stream containing rows of actors and the movies they acted in, along with the director of the particular movie.

Here is the Cypher query to do this. Notice that the direction of the traversal is used to focus on a particular movie during the query:

```
MATCH (a:Person)-[:ACTED_IN]->(m:Movie)<-[:DIRECTED]-(d:Person)  
RETURN a.name, m.title, d.name
```

The result returned is:

a.name	m.title	d.name
"Emil Eifrem"	"The Matrix"	"Lana Wachowski"
"Hugo Weaving"	"The Matrix"	"Lana Wachowski"
"Laurence Fishburne"	"The Matrix"	"Lana Wachowski"
"Carrie-Anne Moss"	"The Matrix"	"Lana Wachowski"
"Keanu Reeves"	"The Matrix"	"Lana Wachowski"
"Emil Eifrem"	"The Matrix"	"Lilly Wachowski"
"Hugo Weaving"	"The Matrix"	"Lilly Wachowski"
"Laurence Fishburne"	"The Matrix"	"Lilly Wachowski"
"Carrie-Anne Moss"	"The Matrix"	"Lilly Wachowski"
"Keanu Reeves"	"The Matrix"	"Lilly Wachowski"
"Hugo Weaving"	"The Matrix Reloaded"	"Lana Wachowski"
"Laurence Fishburne"	"The Matrix Reloaded"	"Lana Wachowski"
"Carrie-Anne Moss"	"The Matrix Reloaded"	"Lana Wachowski"
"Keanu Reeves"	"The Matrix Reloaded"	"Lana Wachowski"

In this query, notice that there are multiple records returned for a movie, each with its set of values for the actor and director.

Later in this training, you will learn other ways to query data and how to control the results returned.

Cypher style recommendations

Here are the **Neo4j-recommended** Cypher coding standards that we use in this training:

- Node labels are CamelCase and begin with an upper-case letter (examples: *Person*, *NetworkAddress*). Note that node labels are case-sensitive.
- Property keys, variables, parameters, and functions are camelCase and begin with a lower-case letter (examples: *businessAddress*, *title*). Note that these elements are case-sensitive.
- Relationship types are in upper-case and can use the underscore.(examples: *ACTED_IN*, *FOLLOWS*). Note that relationship types are case-sensitive and that you **cannot** use the "-" character in a relationship type.
- Cypher keywords are upper-case (examples: **MATCH**, **RETURN**). Note that Cypher keywords are case-insensitive, but a best practice is to use upper-case.
- String constants are in single quotes, unless the string contains a quote or apostrophe (examples: *'The Matrix'*, *"Something's Gotta Give"*).
- Specify variables only when needed for graph processing.
- Place named nodes and relationships (that use variables) before anonymous nodes and relationships in your **MATCH** clauses when possible.
- Specify anonymous relationships with →, ←, or ↔.

Example showing some best coding practices:

```
MATCH (:Person {name: 'Diane Keaton'})-[movRel:ACTED_IN]->(:Movie {title:"Something's Gotta Give"}) RETURN movRel.roles
```

 recommends that you follow the [Cypher Style Guide](#) when writing your Cypher code.

Exercise 3: Filtering queries using relationships

In the query edit pane of Neo4j Browser, execute the browser command: `:play http://guides.neo4j.com/intro-neo4j-exercises` and follow the instructions for Exercise 3.

Check your understanding

Question 1

Suppose you have a graph that contains nodes representing customers and other business entities for your application. The node label in the database for a customer is *Customer*. Each *Customer* node has a property named *email* that contains the customer's email address. What Cypher query do you execute to return the email addresses for all customers in the graph?

Select the correct answer.

- MATCH (n) RETURN n.Customer.email
- MATCH (c:Customer) RETURN c.email
- MATCH (Customer) RETURN email
- MATCH (c) RETURN Customer.email



Question 2

Suppose you have a graph that contains *Customer* and *Product* nodes. A *Customer* node can have a *BOUGHT* relationship with a *Product* node. *Customer* nodes can have other relationships with *Product* nodes. A *Customer* node has a property named *customerName*. A *Product* node has a property named *productName*. What Cypher query do you execute to return all of the products (by name) bought by customer 'ABCCO'.

Select the correct answer.

- MATCH (c:Customer {customerName: 'ABCCO'}) RETURN c.BOUGHT.productName
- MATCH (:Customer 'ABCCO')-[:BOUGHT]->(p:Product) RETURN p.productName
- MATCH (p:Product) -[:BOUGHT_BY]-(:Customer 'ABCCO') RETURN p.productName
- MATCH (:Custom {customerName: 'ABCCO'})-[:BOUGHT]->(p:Product) RETURN p.productName

Question 3

When must you use a variable in a MATCH clause?

Select the correct answer.

- When you want to query the graph using a node label.
- When you specify a property value to match the query.
- When you want to use the node or relationship to return a result.
- When the query involves 2 types of nodes.

Summary

You should now be able to write Cypher code to:

- Retrieve nodes from the graph.
- Filter nodes retrieved using property values of nodes.
- Retrieve property values from nodes in the graph.
- Filter nodes retrieved using relationships.

